

In [209]:

```
1 import pandas as pd
2 import numpy as np
3 import os
4 import re
5 import matplotlib.pyplot as plt
6 import nltk
7 from nltk import RegexpTokenizer
8 from nltk.stem.snowball import SnowballStemmer
9 from sklearn.cluster import KMeans, AgglomerativeClustering, DBSCAN
10 from sklearn.feature_extraction.text import CountVectorizer
11 from sklearn.feature_extraction.text import TfidfVectorizer
12 from wordcloud import WordCloud
13 %matplotlib inline
```

In [2]:

```
1 import warnings
2 warnings.filterwarnings('ignore')
```

## Prepare lyrics data

In [137]:

```
1 # Get lyrics data
2 song_lyrics = pd.read_csv('merge_w_lyrics_full.csv')[['title','artist_x','lyrics']]
3
4 # Clean lyrics
5 song_lyrics['lyrics'] = song_lyrics['lyrics'].str.replace('\n', ' ').replace('.', ' ')
6
7 song_lyrics
```

Out[137]:

	title	artist_x	lyrics	hit
0	10,000 Reasons (Bless the Lord)	Matt Redman	Bless the Lord, O my soul, O my soul Worship H...	0.0
1	100 Proof	Kellie Pickler	Ain't no rain as cold as the look she just gav...	0.0
2	101	Alicia Keys	You used to the sound of a heart that's breaki...	0.0
3	110%	Jessie Ware	Carving my initials on your forehead Now if y...	0.0
4	1313	The Big Pink	Convey your thoughts, translate them well Say ...	0.0
...	...	...	...	...
20977	Danny Phantom	Trippie Redd	Yeah, first I wanna show you some of this dama...	1.0
20978	What's Wrong	Rod Wave	You know what they gon' say When the world tur...	1.0
20979	Demon Time	Trippie Redd	Ha (Ha-ha) Demon time, nigga (Yeah) Gang, uh (...)	1.0
20980	Tick Tock	Young Thug	Yeah, Spider, yeah (Slime) Okay, shit, I just...	1.0
20981	Time Heals	Rod Wave	Ayo Keyz, chop this up, for the one time MarsG...	1.0

20982 rows × 4 columns

In [200]:

```
1 # tokenize and stem lyrics
2 tokenizer = RegexpTokenizer(r'\w+')
3
4 # Instantiate stemmer
5 stemmer = SnowballStemmer('english')
6
7 # List to append stemmed words
8 stemmed = []
9 # List to append tokenized words
10 tokenized = []
11
12 # Create a for loop to iterate through all the rows in specific column
13 for i in song_lyrics['lyrics']:
14
15     # Converting lyrics text to tokens
16     tokens = tokenizer.tokenize(i.lower())
17     tokenized.append(tokens)
18
19     # Stemming all tokens
20     stems = [stemmer.stem(token) for token in tokens]
21     # Appending stems to stemmed list
22     stemmed.append(stems)
23
24 # Creating new dataframe columns
25 song_lyrics['tokenized_lyrics'] = [' '.join(i) for i in tokenized]
26 song_lyrics['stemmed_lyrics'] = [' '.join(i) for i in stemmed]
27
28 song_lyrics
```

Out[200]:

	title	artist_x	lyrics	hit	label_3	tokenized_lyrics	stemmed_lyrics
0	10,000 Reasons (Bless the Lord)	Matt Redman	Bless the Lord, O my soul, O my soul Worship H...	0.0	0	bless the lord o my soul o my soul worship his...	bless the lord o my soul o my soul worship his...
1	100 Proof	Kellie Pickler	Ain't no rain as cold as the look she just gav...	0.0	0	ain t no rain as cold as the look she just gav...	ain t no rain as cold as the look she just gav...
2	101	Alicia Keys	You used to the sound of a heart that's breaki...	0.0	0	you used to the sound of a heart that s breaki...	you use to the sound of a heart that s break i...
3	110%	Jessie Ware	Carving my initials on your forehead Now if y...	0.0	0	carving my initials on your forehead now if yo...	carv my initi on your forehead now if you re n...
4	1313	The Big Pink	Convey your thoughts, translate them well Say ...	0.0	0	convey your thoughts translate them well say t...	convey your thought translat them well say tho...
...	...	...	...	...	...	...	...
20977	Danny Phantom	Trippie Redd	Yeah, first I wanna show you some of this dama...	1.0	1	yeah first i wanna show you some of this damag...	yeah first i wanna show you some of this damag...

	title	artist_x	lyrics	hit	label_3	tokenized_lyrics	stemmed_lyrics
20978	What's Wrong	Rod Wave	You know what they gon' say When the world tur...	1.0	1	you know what they gon say when the world turn...	you know what they gon say when the world turn...
20979	Demon Time	Trippie Redd	Ha (Ha-ha) Demon time, nigga (Yeah) Gang, uh (...)	1.0	1	ha ha ha demon time nigga yeah gang uh gang ga...	ha ha ha demon time nigga yeah gang uh gang ga...
20980	Tick Tock	Young Thug	Yeah, Spider, yeah (Slime) Okay, shit, I just...	1.0	1	yeah spider yeah slime okay shit i just woke u...	yeah spider yeah slime okay shit i just woke u...
20981	Time Heals	Rod Wave	Ayo Keyz, chop this up, for the one time MarsG	1.0	2	ayo keyz chop this up for the one time	ayo keyz chop this up for the one time

In [16]:

```
1 # CountVectorizer
2 cv = CountVectorizer(input='content',
3                      stop_words='english',
4                      lowercase=True,
5                      strip_accents='ascii',
6                      max_features=800)
7
8 dtm_cv = cv.fit_transform(song_lyrics['stemmed_lyrics'])
9
10 df_cv = pd.DataFrame(dtm_cv.toarray(),
11                      columns=cv.get_feature_names(),
12                      index=song_lyrics['title'])
13
14 # Filter non-alphabetic features
15 df_cv = df_cv[[c for c in df_cv.columns if c.isalpha()]]
16
17 df_cv
```

Out[16]:

	abl	abov	act	action	afraid	age	ago	ah	ain	air	...	wrong	ya	ye	yeah	year
title																
10,000 Reasons (Bless the Lord)	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	1
100 Proof	0	0	0	0	0	0	0	0	2	0	...	0	0	0	1	0
101	0	0	0	0	0	0	0	0	0	0	...	0	0	0	1	0
110%	0	0	0	0	0	0	0	0	0	0	...	1	0	0	0	0
1313	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	1
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
Danny Phantom	0	0	0	0	0	0	0	0	0	0	...	0	0	0	32	0
What's Wrong	0	0	1	0	0	0	0	0	3	0	...	2	0	0	23	0
Demon Time	0	0	0	0	0	0	0	0	5	0	...	0	0	0	18	0
Tick Tock	0	0	0	0	0	0	0	0	0	0	...	0	0	0	18	1
Time Heals	0	0	0	0	0	0	0	0	2	0	...	1	0	0	6	0

20982 rows × 793 columns

In [17]:

```
1 # TfidfVectorizer
2 tfidf = TfidfVectorizer(input='content',
3                         stop_words='english',
4                         lowercase=True,
5                         strip_accents='ascii',
6                         max_features=800)
7
8 dtm_tfidf = tfidf.fit_transform(song_lyrics['stemmed_lyrics'])
9
10 df_tfidf = pd.DataFrame(dtm_tfidf.toarray(),
11                         columns=tfidf.get_feature_names(),
12                         index=song_lyrics['title'])
13
14 df_tfidf = df_tfidf[[c for c in df_tfidf.columns if c.isalpha()]]
15
16 df_tfidf
```

Out[17]:

	abl	abov	act	action	afraid	age	ago	ah	ain	air	...	wrong	ya
title													
10,000 Reasons (Bless the Lord)	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.000000	0.0
100 Proof	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.081779	0.0	...	0.000000	0.0
101	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.000000	0.0
110%	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.025611	0.0
1313	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.000000	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
Danny Phantom	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.000000	0.0
What's Wrong	0.0	0.0	0.048145	0.0	0.0	0.0	0.0	0.0	0.085281	0.0	...	0.079602	0.0
Demon Time	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.087408	0.0	...	0.000000	0.0
Tick Tock	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	...	0.000000	0.0
Time Heals	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.060140	0.0	...	0.042102	0.0

20982 rows × 793 columns

In [19]:

```
1 sum(df_tfidf.columns == df_cv.columns)
```

Out[19]:

793

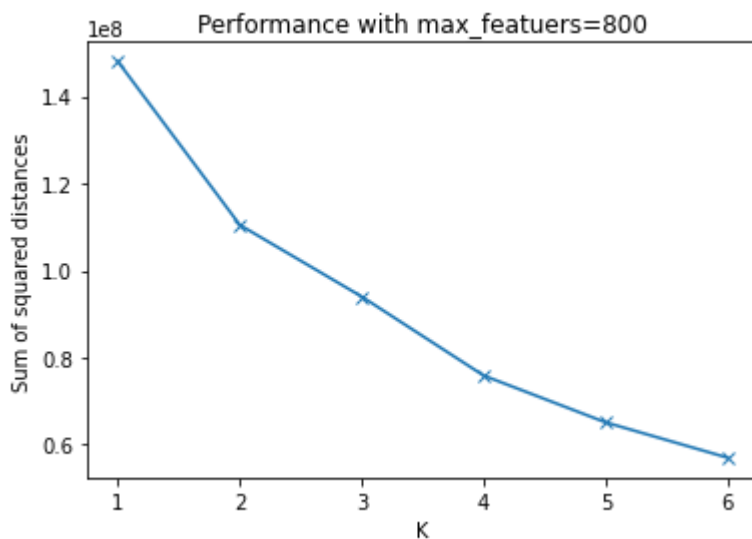
# K-means

In [31]:

```
1 sse = []
2 K = range(1,7)
3
4 for k in K:
5     kmeans = KMeans(n_clusters = k).fit(df_cv)
6     sse.append(kmeans.inertia_)
7
8 plt.plot(K, sse, 'x-')
9 plt.xlabel('K')
10 plt.ylabel('Sum of squared distances')
11 plt.title('Performance with max_featuers=800')
```

Out[31]:

Text(0.5, 1.0, 'Performance with max\_featuers=800')



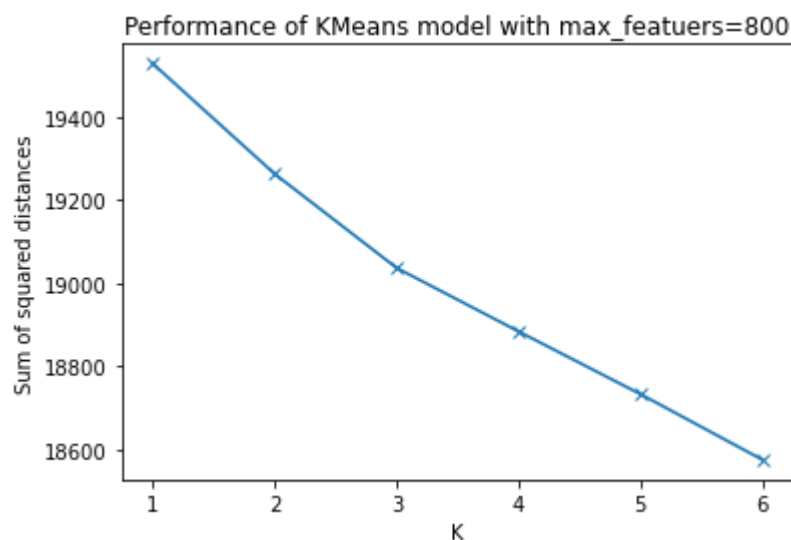
The plot tells a lot. With some pre-knowledge, the lyrics belong two major categories - hit song and non-hit song. Therefore, there is an elbow point at  $k = 2$ . But hit and non-hit seems not sufficient.  $k = 4$  is also somehow reasonable because it is possible that different genre may produce different styles of lyrics.

In [28]:

```
1 # Use TfidfVectorizer
2 sse = []
3 K = range(1,7)
4
5 for k in K:
6     kmeans = KMeans(n_clusters = k).fit(df_tfidf)
7     sse.append(kmeans.inertia_)
8
9 plt.plot(K, sse, 'x-')
10 plt.xlabel('K')
11 plt.ylabel('Sum of squared distances')
12 plt.title('Performance of KMeans model with max_featuers=800')
```

Out[28]:

Text(0.5, 1.0, 'Performance of KMeans model with max\_featuers=800')



Using TfidfVectorizer,  $k = 3$  gives a clear elbow. This makes sense because there might be songs pretty popular but do not feature on Billboard Hot 100 charts.

### **$k=2$ with CounterVectorizer**

In [59]:

```
1 model_2 = KMeans(n_clusters = 2)
2 model_2.fit(df_cv)
3
4 # accuracy
5 sum(model_2.labels_ == song_lyrics['hit'])/len(song_lyrics)
```

Out[59]:


0.8279001048517777

Cluster visualization using TruncatedSVD



In [218]:

```
1 from sklearn.decomposition import TruncatedSVD
2
3 svd = TruncatedSVD(n_components=3, random_state=42)
4 dtm_cv_3d = svd.fit_transform(dtm_cv)
5
6 plt.figure(figsize=(6,4))
7 plt.scatter(dtm_cv_3d[:, 1], dtm_cv_3d[:, 2], s=0.8, c=model_2.labels_)
8 plt.xlabel('dim_1')
9 plt.ylabel('dim_2')
10 plt.xlim(0,800)
11 plt.ylim(-10,10)
12 plt.title('Clusters of KMeans(2) with CV')
13 plt.show()
```



In [183]:


```
1 model_2_t = KMeans(n_clusters = 2)
2 model_2_t.fit(df_cv)
3
4 # accuracy
5 sum(model_2_t.labels_ == song_lyrics['hit'])/len(song_lyrics)
```

Out[183]:

0.8279001048517777

In [213]:

```
1 dtm_tfidf_3d = svd.fit_transform(dtm_tfidf)
2 plt.figure(figsize=(6,4))
3 plt.scatter(dtm_tfidf_3d[:, 1], dtm_tfidf_3d[:, 2], s=0.8, c=model_2_t.labels_)
4 plt.xlabel('dim_1')
5 plt.ylabel('dim_2')
6 plt.title('Clusters of KMeans(2) with Tfidf')
7 plt.show()
```



The accuracy of KMeans clustering with  $k = 2$  is around 83%, which is not bad (because we cannot tell hit songs merely by the lyrics). The visualization of clusters does not reveal too much information when  $k = 2$ , as most of the songs are non-hit songs and only 3300 are hit songs; therefore, the majorities are purple points (non-hit songs). Specifically, when comparing CountVectorizer and TfidfVectorizer, the clustering results given by TfidfVectorizer is much more satisfying than CountVectorizer, as it is more centered.

Let's try more number of clusters.

**$k=3$**

In [155]:

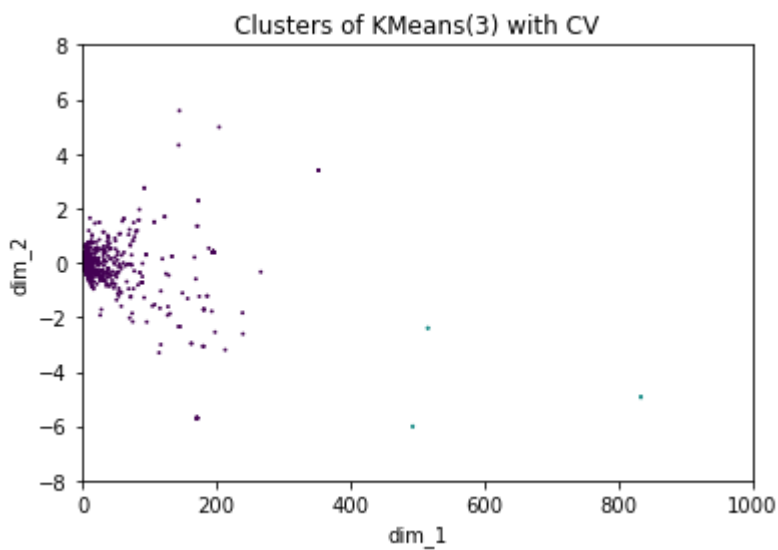
```
1 # CountVectorizer
2 model_3 = KMeans(n_clusters = 3)
3 model_3.fit(df_cv)
4
5 # TfidfVectorizer
6 model_3_t = KMeans(n_clusters = 3)
7 model_3_t.fit(df_tfidf)
```

Out[155]:

KMeans(n\_clusters=3)

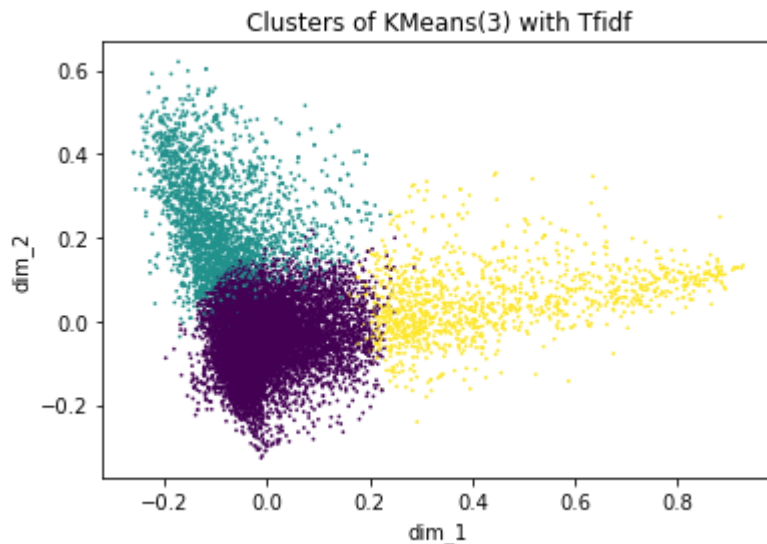
In [176]:

```
1 plt.figure(figsize=(6,4))
2 plt.scatter(dtm_cv_3d[:, 1], dtm_cv_3d[:, 2], s=0.8, c=model_3.labels_)
3 plt.xlabel('dim_1')
4 plt.ylabel('dim_2')
5 plt.xlim(-1,1000)
6 plt.ylim(-8,8)
7 plt.title('Clusters of KMeans(3) with CV')
8 plt.show()
```



In [215]:

```
1 plt.figure(figsize=(6,4))
2 plt.scatter(dtm_tfidf_3d[:, 1], dtm_tfidf_3d[:, 2], s=0.6, c=model_3_t.labels_)
3 plt.xlabel('dim_1')
4 plt.ylabel('dim_2')
5 plt.title('Clusters of KMeans(3) with Tfidf')
6 plt.show()
```



Aligned with what is observed using  $k = 2$ , the clusters obtained using TfidfVectorizer are much more reliable than CountVectorizer when  $k = 3$ . It can be easily seen that  $k = 3$  is a very good choice of number of clusters because cases in each group are tightly clustered without much overlapping.

**$k=5$**

In [164]:


```
1 # CountVectorizer
2 model_5 = KMeans(n_clusters = 5)
3 model_5.fit(df_cv)
4
5 # TfidfVectorizer
6 model_5_t = KMeans(n_clusters = 5)
7 model_5_t.fit(df_tfidf)
```

Out[164]:

KMeans(n\_clusters=5)

In [216]:

```
1 plt.figure(figsize=(6,4))
2 plt.scatter(dtm_tfidf_3d[:, 1], dtm_tfidf_3d[:, 2], s=0.6, c=model_5_t.labels_)
3 plt.xlabel('dim_1')
4 plt.ylabel('dim_2')
5 plt.title('Clusters of KMeans(5) with CV')
6 plt.show()
```



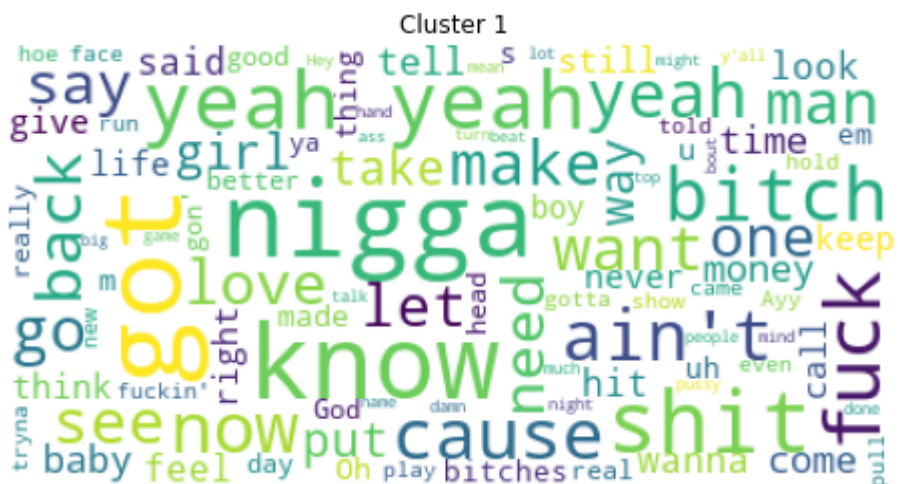
$k = 5$  might not be a good idea, since the separation does not seem maximized as  $k = 3$ .

## Wordcloud of each cluster

$k = 3$

In [206]:


```
1 song_lyrics['label_3'] = model_3_t.labels_
2 for i in range(3):
3     text = ' '.join(song_lyrics.loc[song_lyrics['label_3']==i, 'lyrics'])
4     plt.figure(figsize=(8,6))
5     wordcloud = WordCloud(max_font_size=50, max_words=100, background_color='white')
6     plt.imshow(wordcloud)
7     plt.axis('off')
8     plt.title('Cluster %s' %i)
9     plt.show()
```



## Hierarchical clustering

In [272]:

```
1 import scipy.cluster.hierarchy as sch
2
3 plt.figure(figsize=(20,8))
4 dendrogram = sch.dendrogram(sch.linkage(df_tfidf, method = 'ward'))
5 plt.title('Dendrogram')
6 plt.ylabel('Euclidean distances')
7 plt.show()
```



In [229]:

```
1 hc_euc = AgglomerativeClustering(n_clusters=3,
2                                   affinity = 'euclidean',
3                                   linkage = 'single')
4
5 hc_euc.fit(df_tfidf)
```

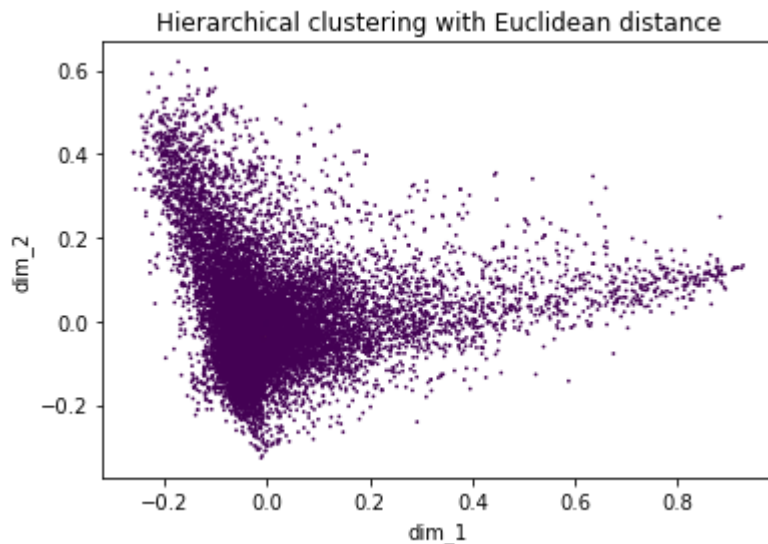
Out[229]:

AgglomerativeClustering(linkage='single', n\_clusters=3)




In [241]:

```
1 plt.figure(figsize=(6,4))
2 plt.scatter(dtm_tfidf_3d[:, 1], dtm_tfidf_3d[:, 2], s=0.6, c=hc_euc.labels_)
3 plt.xlabel('dim_1')
4 plt.ylabel('dim_2')
5 plt.title('Hierarchical clustering with Euclidean distance')
6 plt.show()
```



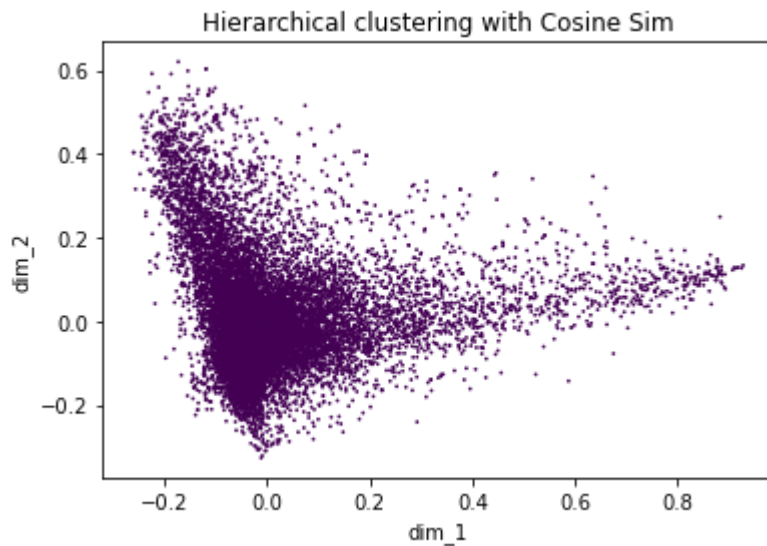
In [239]:

```
1 hc_man = AgglomerativeClustering(n_clusters=3,
2                                 affinity = 'manhattan',
3                                 linkage = 'single')
4
5 hc_man.fit(df_tfidf)
6
7 plt.figure(figsize=(6,4))
8 plt.scatter(dtm_tfidf_3d[:, 1], dtm_tfidf_3d[:, 2], s=0.6, c=hc_man.labels_)
9 plt.xlabel('dim_1')
10 plt.ylabel('dim_2')
11 plt.title('Hierarchical clustering with Manhattan distance')
12 plt.show()
```



In [240]:

```
1 hc_cos = AgglomerativeClustering(n_clusters=3,  
2                                 affinity = 'cosine',  
3                                 linkage = 'single')  
4  
5 hc_cos.fit(df_tfidf)  
6  
7 plt.figure(figsize=(6,4))  
8 plt.scatter(dtm_tfidf_3d[:, 1], dtm_tfidf_3d[:, 2], s=0.6, c=hc_cos.labels_)  
9 plt.xlabel('dim_1')  
10 plt.ylabel('dim_2')  
11 plt.title('Hierarchical clustering with Cosine Sim')  
12 plt.show()
```



## DBSCAN


Determine epsilon

In [256]:

```
1 from sklearn.neighbors import NearestNeighbors
2
3 nbrs = NearestNeighbors(n_neighbors=3).fit(dtm_tfidf)
4 distances, indices = nbrs.kneighbors(dtm_tfidf)
5 distances = np.sort(distances, axis=0)
6 distances = distances[:,1]
7 plt.plot(distances)
```

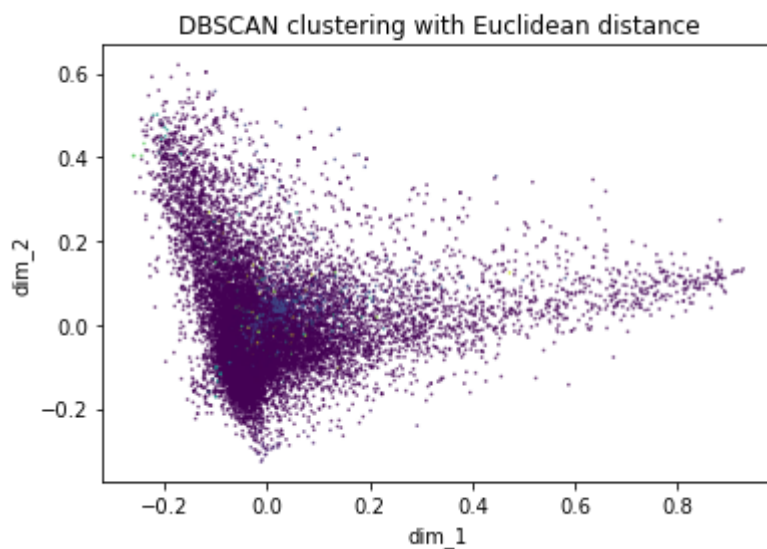
Out[256]:

[matplotlib.lines.Line2D at 0x7fb5cc747820>]




In [269]:

```
1 db_euc = DBSCAN(eps=0.62, metric='euclidean', n_jobs=-1)
2 db_euc.fit(dtm_tfidf)
3
4 plt.figure(figsize=(6,4))
5 plt.scatter(dtm_tfidf_3d[:, 1], dtm_tfidf_3d[:, 2], s=0.6, c=db_euc.labels_, alp
6 plt.xlabel('dim_1')
7 plt.ylabel('dim_2')
8 plt.title('DBSCAN clustering with Euclidean distance')
9 plt.show()
```



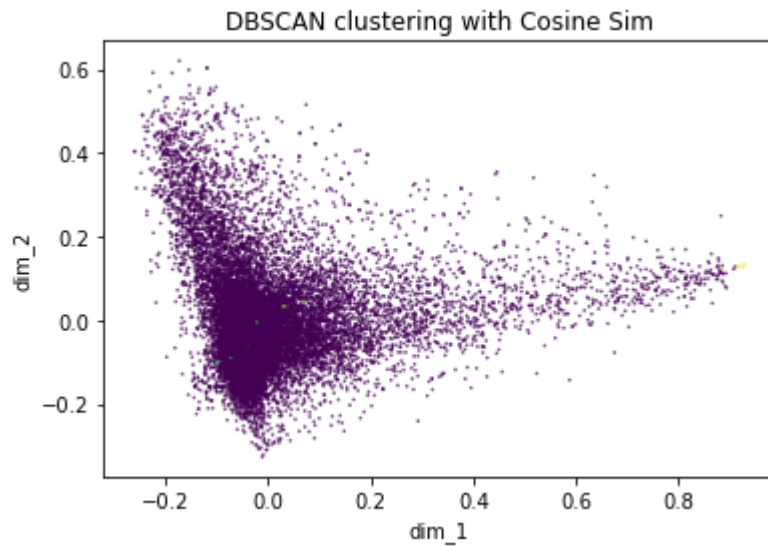
In [270]:

```
1 db_man = DBSCAN(eps=0.62, metric='manhattan', n_jobs=-1)
2 db_man.fit(dtm_tfidf)
3
4 plt.figure(figsize=(6,4))
5 plt.scatter(dtm_tfidf_3d[:, 1], dtm_tfidf_3d[:, 2], s=0.6, c=db_man.labels_, alp
6 plt.xlabel('dim_1')
7 plt.ylabel('dim_2')
8 plt.title('DBSCAN clustering with Manhattan distance')
9 plt.show()
```



In [271]:

```
1 db_cos = DBSCAN(eps=0.62, metric='manhattan', n_jobs=-1)
2 db_cos.fit(dtm_tfidf)
3
4 plt.figure(figsize=(6,4))
5 plt.scatter(dtm_tfidf_3d[:, 1], dtm_tfidf_3d[:, 2], s=0.6, c=db_cos.labels_, alp
6 plt.xlabel('dim_1')
7 plt.ylabel('dim_2')
8 plt.title('DBSCAN clustering with Cosine Sim')
9 plt.show()
```



In [ ]:

1